

Reinforce Explanation for LUTEnv

Step 1: You start at the initial netlist state




The netlist as read from the .BLIF files

Step 2: Each different move puts you in a different netlist state

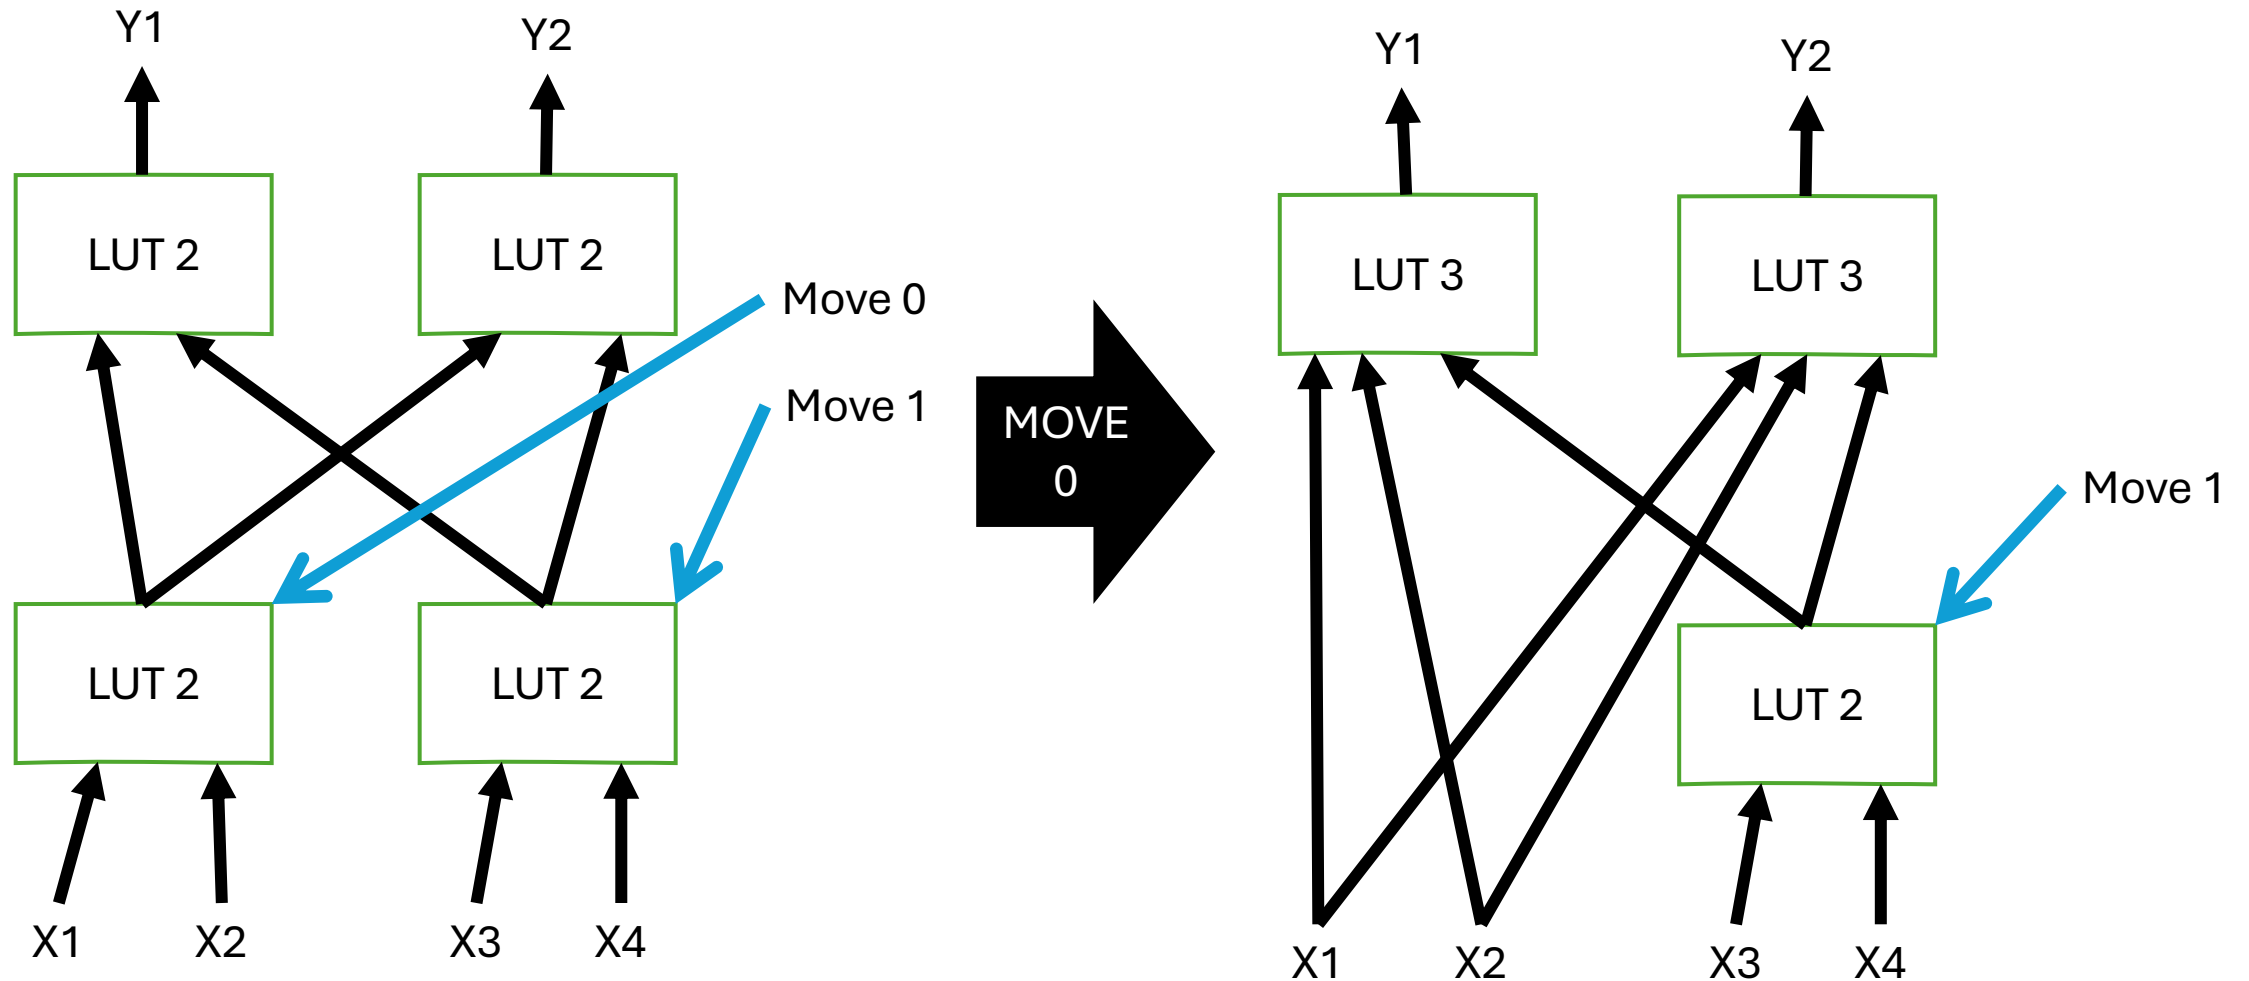
Reminder:

Move = LUT Merges

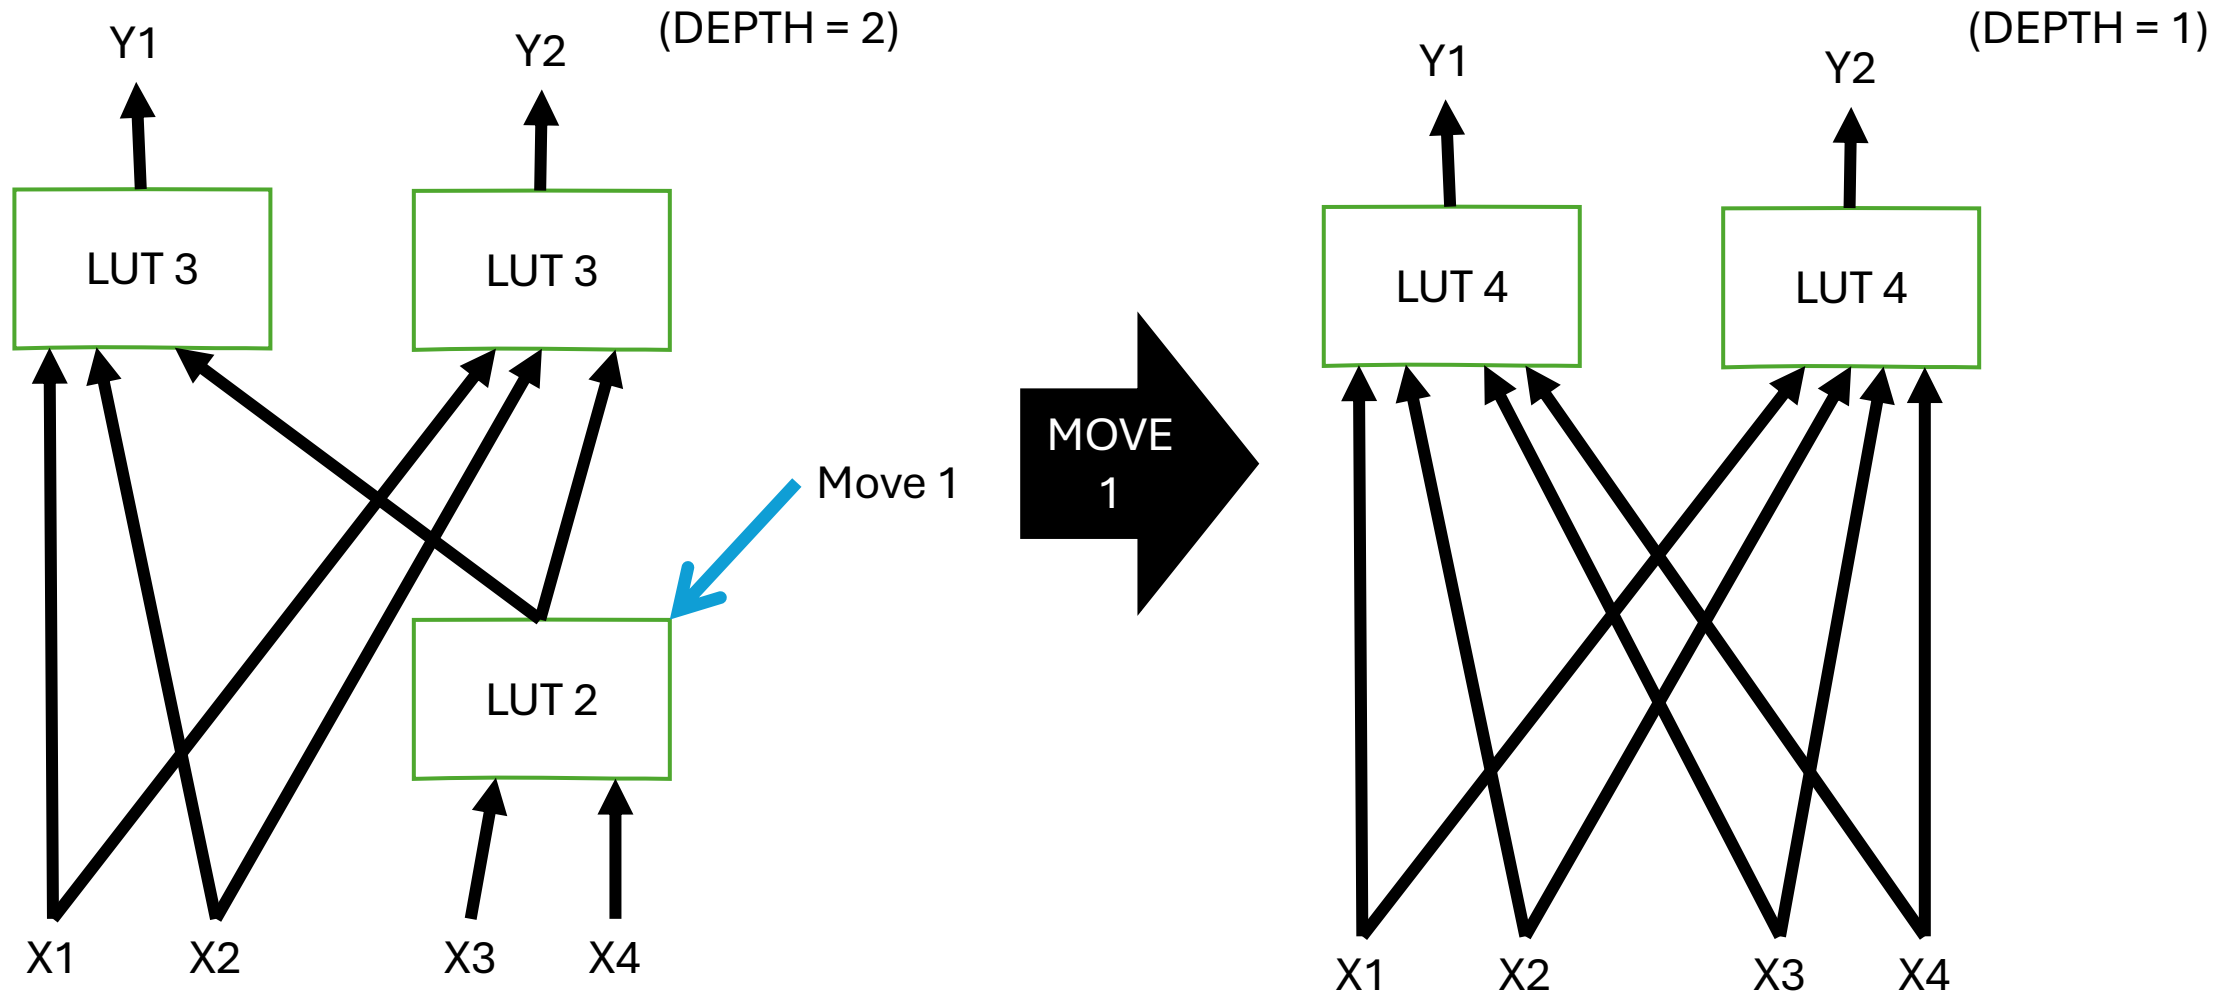
Which change the netlist
(new state)



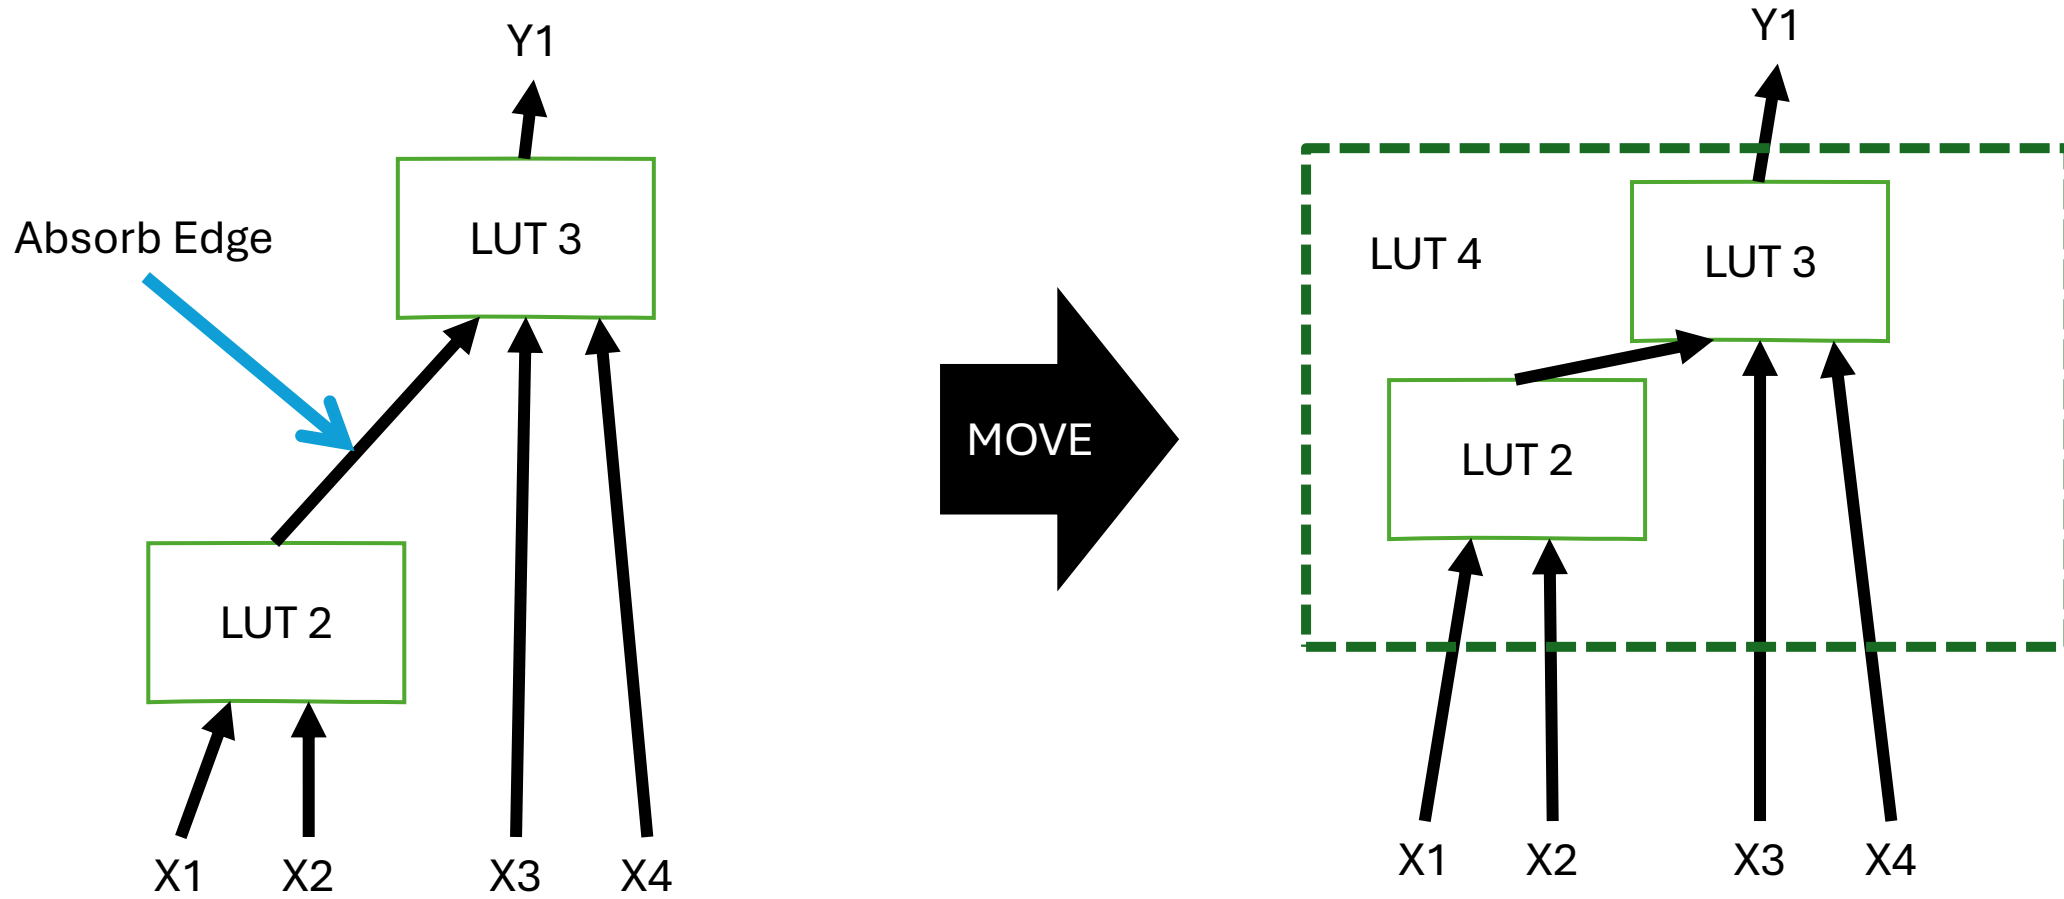
Concrete Representation Of Moves




Concrete Representation Of Moves



How Functional Equivalence is Maintained



In general, this forms a Directed Acyclic Graph



Directed because
Undo moves aren't
provided

Graph and not tree
because a state can be
reached in multiple
move orders

Some states have no moves

All remaining moves (LUT merges) create LUTs with $K > 6$ and violate constraints.


Therefore, the state is terminal

Terminal States have some given reward / value
 $\text{Value} = -\text{Depth} * \text{\#LUTs}$



A policy function will suggest moves for a state (Vector of probabilities)

$$\pi(s_x, \theta) = [p_A, p_B, p_C]$$



Policy gradients is an update to the weights

$$\nabla_{\theta} = \nabla_{\theta}^w \log(\pi(s_x, \theta))$$


Is an update to the weights that makes

p_w (probability move is move w) more likely (in state s_x)
probability move is everything else less likely (in state s_x)
When added to θ

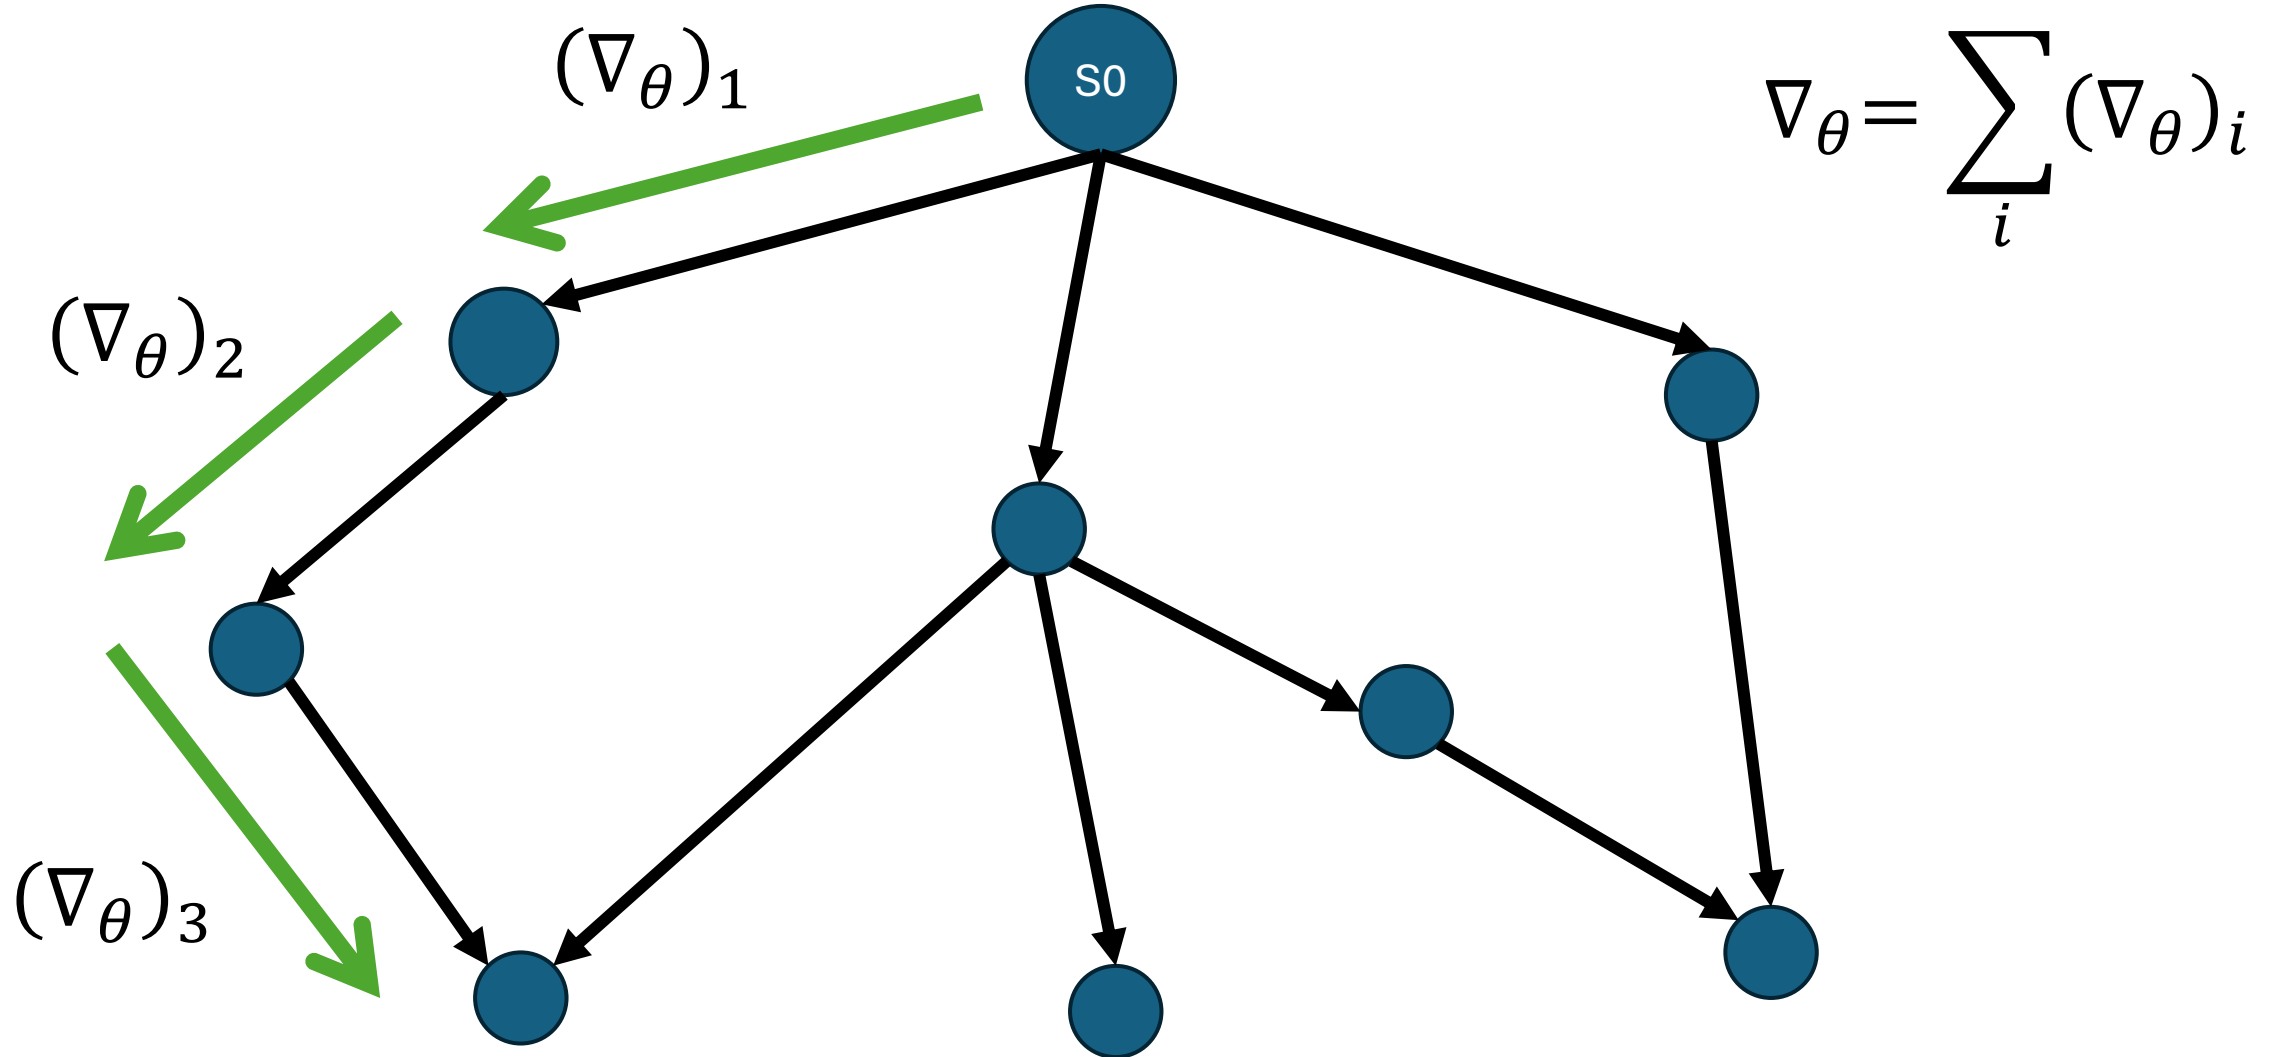
The gradient is taking w.r.t log probability for mathematical reasons.

$\theta_{after} = \theta_{before} + \nabla_{\theta}$: Move w is made more likely in state s_x

During an episode you take some fixed number of random steps (or until you reach some terminal state for full rollout)



Accumulate the gradient over those steps



Consider the rewards of that episode

R = Rewards earned in episode (- Final LUTs * Final depth)

B = Baseline rewards earned by policy (Assume you have some way of computing how much reward is “typically earned”)

$A = \text{Advantage} = R - B$

$A > 0$: The moves were better than expected

$A < 0$: The moves were worse than expected

Update the policy based on advantage

$$\theta_{final} = \theta_{initial} + A * \nabla_{\theta}$$

If $A > 0$: All moves made are updated to be made more likely.

If $A < 0$: All move made are updated to be made less likely.

You might use ADAMW or other optimizers, this is simple gradient ascent.

Baselines

Baseline can be done in many ways:

- Training and managing a value function to estimate baseline
- Population of episodes and computing mean and std deviation of rewards in the population.
- (In this lab) exponential moving averages of rewards earned in previous episodes.

Policy Gradient Methods are All variants of this Formulation

Variants can be quite sophisticated:

- Dynamic policy updates during an episode
- Off – policy existing reward data (Difficult to do in practice)
- Managing step sizes and proximity to original policy
- Managing reward credit assignment to particular moves
- ...